

Intro to pandas data structures

October 26, 2013 / Tags: python pandas sql tutorial data science

UPDATE: If you're interested in learning pandas from a SQL perspective and would prefer to watch a video, you can find video of my 2014 PyData NYC talk [here](#).



However, the other week a couple of coworkers expressed their interest in learning a bit more about it - this seemed like a good reason to revisit the topic.

What follows is a fairly thorough introduction to the library. I chose to break it into three parts as I felt it was too long and daunting as one.

- Part 1: Intro to pandas data structures, covers the basics of the library's two main data structures - Series and DataFrames.
- Part 2: Working with DataFrames, dives a bit deeper into the functionality of DataFrames. It shows how to inspect, select, filter, merge, combine, and group your data.
- Part 3: Using pandas with the MovieLens dataset, applies the learnings of the first two parts in order to answer a few basic analysis questions about the MovieLens ratings data.

If you'd like to follow along, you can find the necessary CSV files [here](#) and the MovieLens dataset [here](#).

**Greg Reda**  

BLOG

in   

...ing, for data cleaning, etc. It has always been great for prepping and munging data, but it's never been great for analysis - you'd usually end up using R or loading it into a database and using SQL (or worse, Excel). pandas makes Python great for analysis.

Data Structures

pandas introduces two new data structures to Python - Series and DataFrame, both of which are built on top of NumPy (this means it's fast).

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
pd.set_option('max_columns', 50)
%matplotlib inline
```

Series

A Series is a one-dimensional object similar to an array, list, or column in a table. It will assign a labeled index to each item in the Series. By default, each item will receive an index label from 0 to N, where N is the length of the Series minus one.

Greg Reda  

BLOG

in G+  

```

0          7
1      Heisenberg
2          3.14
3      -1789710578
4      Happy Eating!
dtype: object

```

Alternatively, you can specify an index to use when creating the Series.

```

s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'],
              index=['A', 'Z', 'C', 'Y', 'E'])
s

```

```

A          7
Z      Heisenberg
C          3.14
Y      -1789710578
E      Happy Eating!
dtype: object

```

The Series constructor can convert a dictionary as well, using the keys of the dictionary as its index.

```

d = {'Chicago': 1000, 'New York': 1300, 'Portland': 900,
     'Austin': 450, 'Boston': None}
cities = pd.Series(d)
cities

```

```

Austin          450
Boston          NaN
Chicago         1000
New York        1300
Portland         900
San Francisco   1100
dtype: float64

```

You can use the index to select specific items from the Series ...

```

cities['Chicago']

```

```

1000.0

```

**Greg Reda**  

BLOG

in   

```
Chicago      1000
Portland     900
San Francisco 1100
dtype: float64
```

Or you can use boolean indexing for selection.

```
cities[cities < 1000]
```

```
Austin      450
Portland    900
dtype: float64
```

That last one might be a little weird, so let's make it more clear - `cities < 1000` returns a Series of True/False values, which we then pass to our Series `cities`, returning the corresponding True items.

```
less_than_1000 = cities < 1000
print(less_than_1000)
print('\n')
print(cities[less_than_1000])
```

```
Austin      True
Boston      False
Chicago     False
New York    False
Portland     True
San Francisco False
dtype: bool
```

```
Austin      450
Portland    900
dtype: float64
```

You can also change the values in a Series on the fly.

**Greg Reda**  

BLOG

in G+  

```
('Old value:', 1000.0)
('New value:', 1400.0)
```

```
# changing values using boolean logic
print(cities[cities < 1000])
print('\n')
cities[cities < 1000] = 750

print cities[cities < 1000]
```

```
Austin      450
Portland    900
dtype: float64
```

```
Austin      750
Portland    750
dtype: float64
```

What if you aren't sure whether an item is in the Series? You can check using idiomatic Python.

```
print('Seattle' in cities)
print('San Francisco' in cities)
```

```
False
True
```

Mathematical operations can be done using scalars and functions.

```
# divide city values by 3
cities / 3
```

```
Austin      250.000000
Boston              NaN
Chicago      466.666667
New York     433.333333
Portland     250.000000
San Francisco 366.666667
dtype: float64
```

**Greg Reda**  

BLOG

in   

Austin	562500
Boston	NaN
Chicago	1960000
New York	1690000
Portland	562500
San Francisco	1210000

dtype: float64

You can add two Series together, which returns a union of the two Series with the addition occurring on the shared index values. Values on either Series that did not have a shared index will produce a NULL/NaN (not a number).

```
print(cities[['Chicago', 'New York', 'Portland'])
print('\n')
print(cities[['Austin', 'New York']])
print('\n')
print(cities[['Chicago', 'New York', 'Portland']])
```

Chicago	1400
New York	1300
Portland	750

dtype: float64

Austin	750
New York	1300

dtype: float64

Austin	NaN
Chicago	NaN
New York	2600
Portland	NaN

dtype: float64



Greg Reda  

BLOG

in G+  

NULL/NaN values.

NULL checking can be performed with
isnull and notnull.

```
# returns a boolean series indicating which values are not null
cities.notnull()
```

```
Austin      True
Boston      False
Chicago     True
New York    True
Portland    True
San Francisco True
dtype: bool
```

```
# use boolean logic to grab the NULL cities
print(cities.isnull())
print('\n')
print(cities[cities.isnull()])
```

```
Austin      False
Boston      True
Chicago     False
New York    False
Portland    False
San Francisco False
dtype: bool
```

```
Boston      NaN
dtype: float64
```


**Greg Reda**  

BLOG

in   

structure comprised of rows and columns, akin to a spreadsheet, database table, or R's `data.frame` object. You can also think of a `DataFrame` as a group of `Series` objects that share an index (the column names).

For the rest of the tutorial, we'll be primarily working with `DataFrames`.

Reading Data

To create a `DataFrame` out of common Python data structures, we can pass a dictionary of lists to the `DataFrame` constructor.

Using the `columns` parameter allows us to tell the constructor how we'd like the columns ordered. By default, the `DataFrame` constructor will order the columns alphabetically (though this isn't the case when reading from a file - more on that next).

**Greg Reda**  

BLOG

in   

	year	team	wins	losses
0	2010	Bears	11	5
1	2011	Bears	8	8
2	2012	Bears	10	6
3	2011	Packers	15	1
4	2012	Packers	11	5
5	2010	Lions	6	10
6	2011	Lions	10	6
7	2012	Lions	4	12

Much more often, you'll have a dataset you want to read into a DataFrame. Let's go through several common ways of doing so.

CSV

Reading a CSV is as simple as calling the `read_csv` function. By default, the `read_csv` function expects the column separator to be a comma, but you can change that using the `sep` parameter.

```
%cd ~/Dropbox/tutorials/pandas/  
/Users/gjreda/Dropbox (Personal)/t
```

**Greg Reda**  

BLOG

in   

```
Year, Age, Tm, Lg, W, L, W-L%, ERA, G, GS, GF
1995, 25, NYY, AL, 5, 3, .625, 5.51, 19, 10, 1
1996, 26, NYY, AL, 8, 3, .727, 2.09, 61, 0, 1
1997, 27, NYY, AL, 6, 4, .600, 1.88, 66, 0, 5
1998, 28, NYY, AL, 3, 0, 1.000, 1.91, 54, 0, 1
```

```
from_csv = pd.read_csv('mariano-rivera.csv')
from_csv.head()
```

	Year	Age	Tm	Lg	W	L	W-L%
0	1995	25	NYN	AL	5	3	0.625
1	1996	26	NYN	AL	8	3	0.727
2	1997	27	NYN	AL	6	4	0.600
3	1998	28	NYN	AL	3	0	1.000
4	1999	29	NYN	AL	4	3	0.571

Our file had headers, which the function inferred upon reading in the file. Had we wanted to be more explicit, we could have passed `header=None` to the function along with a list of column names to use:

```
# Source: pro-football-reference.com/players/M/M
!head -n 5 peyton-passing-TDs-2012.csv
```

```
1,1,2012-09-09,DEN,,PIT,W 31-19,3,1
2,1,2012-09-09,DEN,,PIT,W 31-19,4,1
3,2,2012-09-17,DEN,@,ATL,L 21-27,2,1
4,3,2012-09-23,DEN,,HOU,L 25-31,4,3
5,3,2012-09-23,DEN,,HOU,L 25-31,4,6
```



	num	game	date	team
0	1	1	2012-09-09	DEN
1	2	1	2012-09-09	DEN
2	3	2	2012-09-17	DEN
3	4	3	2012-09-23	DEN
4	5	3	2012-09-23	DEN

pandas' various *reader* functions have many parameters allowing you to do things like skipping lines of the file, parsing dates, or specifying how to handle NA/NULL datapoints.

There's also a set of *writer* functions for writing to a variety of formats (CSVs, HTML tables, JSON). They function exactly as you'd expect and are typically called `to_format`:

```
my_dataframe.to_csv('path_to_file.csv')
```

Take a look at the IO documentation to familiarize yourself with file reading/writing functionality.

**Greg Reda**  

BLOG

in   

do, too. Thankfully, pandas allows you to read and write Excel files, so you can easily read from Excel, write your code in Python, and then write back out to Excel - no need for VBA.

Reading Excel files requires the xlr library. You can install it via pip (*pip install xlr*).

Let's first write a DataFrame to Excel.

```
# this is the DataFrame we created from a dictio
football.head()
```

	year	team	wins	losses
0	2010	Bears	11	5
1	2011	Bears	8	8
2	2012	Bears	10	6
3	2011	Packers	15	1
4	2012	Packers	11	5

```
# since our index on the football DataFrame is m
football.to_excel('football.xlsx', index=False)
```

```
!ls -l *.xlsx
```

```
-rw-r--r--@ 1 gjreda  staff  5665 1
```

```
# delete the DataFrame
del football
```

**Greg Reda**  

BLOG

in   

	year	team	wins	losses
0	2010	Bears	11	5
1	2011	Bears	8	8
2	2012	Bears	10	6
3	2011	Packers	15	1
4	2012	Packers	11	5
5	2010	Lions	6	10
6	2011	Lions	10	6
7	2012	Lions	4	12

Database

pandas also has some support for reading/writing DataFrames directly from/to a database [docs]. You'll typically just need to pass a connection object or sqlalchemy engine to the `read_sql` or `to_sql` functions within the `pandas.io` module.

Note that `to_sql` executes as a series of INSERT INTO statements and thus trades speed for simplicity. If you're writing a large DataFrame to a database, it might be quicker to write the DataFrame to CSV and load that directly using the database's file import arguments.

**Greg Reda**  

BLOG

in   

```
results = sql.read_sql(query, con=conn)
results.head()
```

	tow_date	make	style	model
0	01/19/2013	FORD	LL	
1	01/19/2013	FORD	4D	
2	01/19/2013	FORD	4D	
3	01/19/2013	FORD	LL	
4	01/19/2013	FORD	LL	



Greg Reda  

BLOG

in G+  

read directly into a DataFrame, I prefer to read the results directly from the clipboard. I'm often tweaking queries in my SQL client (Sequel Pro), so I would rather see the results *before* I read it into pandas. Once I'm confident I have the data I want, then I'll read it into a DataFrame.

This works just as well with any type of delimited data you've copied to your clipboard. The function does a good job of inferring the delimiter, but you can also use the sep parameter to be explicit.

Hank Aaron

Standard Batting		More Stats		Glossary · Show Minors Stats · SHARE · Embed · CSV · PSE · LINK · 7 · Advanced Search																									
Year	Age	Tm	Lg	PA	AB	R	H	2B	3B	HR	RBI	SB	CS	BB	SO	BA	OBP	SLG	OPS	OPS+	TS	GDP	HBP	SH	SF	LOB	Pct	Awards	
1954	20	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1955	21	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1956	22	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1957	23	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1958	24	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1959	25	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1960	26	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1961	27	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1962	28	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1963	29	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1964	30	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1965	31	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1966	32	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1967	33	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1968	34	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1969	35	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1970	36	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1971	37	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1972	38	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1973	39	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1974	40	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1975	41	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
1976	42	SEA	NA	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	

**Greg Reda**  

BLOG

in   

	Year	Age	Tm	Lg	G	PA
0	1954	20	MLN	NL	122	509
1	1955 ★	21	MLN	NL	153	665
2	1956 ★	22	MLN	NL	153	660
3	1957 ★	23	MLN	NL	151	675
4	1958 ★	24	MLN	NL	153	664

URL

With `read_table`, we can also read directly from a URL.

Let's use the best sandwiches data that I wrote about scraping a while back.

**Greg Reda** [Twitter](#) [GitHub](#)

BLOG

[in](#) [G+](#) [RSS](#)

	rank	sandwich	restaurant
0	1	BLT	Old Oak Tap
1	2	Fried Bologna	Au Cheval
2	3	Woodland Mushroom	Xoco

Google Analytics

pandas also has some integration with the Google Analytics API, though there is some setup required. I won't be covering it, but you can read more about it [here](#) and [here](#).

Move onto the next section, which covers working with DataFrames.

[Tweet](#)[Share](#) 49[Vote](#)

124

[← Back to Home](#)



Greg Reda  

[BLOG](#)

[in](#) [G+](#)  
